

Secure Software Design and Development



Final Report

Submitted to:

Engr Muhammad Ahmad Nawaz

Prepared by:

Maheen Aamir (SP23-BCT-027)

Malik Muhammad Huzaifa (SP23-BCT-029)

Muzamil Ali Akhtar (SP23-BCT-043)

Maryum Masood (SP23-BCT-031)

Department of computer science

(Cyber Security)

COMSATS UNIVERSITY ISLAMABAD

Contents

Abstract	3
I. INTRODUCTION	3
A. Project Objectives	4
B. Document Organization	4
II. PROTECTION NEEDS ELICITATION (PNE)	4
A. Asset Identification and Classification	4
B. Threat Context Analysis	5
C. Elicited Security Requirements	5
III. THREAT MODELING AND ATTACK SURFACE ANALYSIS	5
A. STRIDE Summary Analysis	6
B. Data Flow Diagram (DFD) Trust Boundary Analysis	6
C. Attack Tree Decomposition	7
IV. CVSS v3.1 RISK ASSESSMENT	8
A. IDOR — Detailed CVSS Analysis (Base Score: 6.5 Medium)	8
B. CSRF — Detailed CVSS Analysis (Base Score: 8.8 High)	8
C. Clickjacking — Detailed CVSS Analysis (Base Score: 5.4 Medium)	8
V. VULNERABILITY FINDINGS AND REMEDIATION REQUIREMENTS	9
A. Finding 1: Insecure Direct Object Reference (IDOR)	9
B. Finding 2: Cross-Site Request Forgery (CSRF)	9
C. Finding 3: Clickjacking (UI Redress Attack)	9
VI. SECURITY IMPLEMENTATION (CODE-LEVEL)	9
A. IDOR Mitigation — User-Scoped Object Retrieval	10
B. CSRF Mitigation — Synchronizer Token Pattern	10
C. Session Hardening and Clickjacking Prevention	11
D. Automated Security Test Suite	11
VII. CI/CD PIPELINE INTEGRATION AND DEVSECOPS AUTOMATION	12
A. Trigger Strategy	12
B. Pipeline Job Architecture	12
C. GitHub Actions Workflow Configuration	13
D. Artifact Generation and Evidence Retention	14
VIII. SAST + DAST RESULTS AND VALIDATION	14
A. SAST Results — CodeQL Analysis	14
B. DAST Results — OWASP ZAP Baseline Scan	15
C. Unit Test Execution Evidence	15
IX. CONCLUSION AND FUTURE HARDENING RECOMMENDATIONS	16
A. Summary of Achievements	16
B. Residual Risks and Future Hardening Roadmap	16
C. DevSecOps Maturity Assessment	17

X. ACKNOWLEDGMENTS	17
APPENDICES	17
Appendix A — Evidence Checklist	17
Appendix B — Full STRIDE Threat Matrix (Extended)	18
Appendix C — CVSS v3.1 Scoring Breakdown	18
Appendix D — CI/CD Workflow YAML (Full)	19
Appendix E — Security Unit Test Suite (test_security.py)	19
Appendix F — Data Flow Diagram	19

Abstract

Modern web applications are increasingly targeted through vulnerabilities arising from improper access control, object-level authorization flaws, and weak request integrity rather than cryptographic failures. This report presents a comprehensive security hardening workflow applied to a Flask-based To-Do web application within a 48-hour DevSecOps sprint, simulating the time-constrained environment of a real-world secure software development lifecycle. The project focused on identifying, analyzing, and mitigating critical web vulnerabilities aligned with the OWASP Top 10, specifically targeting Insecure Direct Object Reference (IDOR), Cross-Site Request Forgery (CSRF), and Clickjacking.

The methodology employed a structured multi-step approach. Protection Needs Elicitation (PNE) was conducted to identify and classify assets according to confidentiality, integrity, and availability (CIA) requirements. STRIDE-based threat modeling decomposed the system's attack surface, mapping each potential adversarial action to a threat category, while CVSS v3.1 scoring quantified risk severity and informed remediation priorities. Security measures were implemented at multiple layers: route logic, Jinja2 templates, session management, and server configuration, ensuring proper object-level authorization, CSRF validation, and anti-clickjacking headers.

A DevSecOps pipeline using GitHub Actions was developed to enforce shift-left security principles, integrating automated unit tests, static application security testing (SAST) via CodeQL, dynamic application security testing (DAST) via OWASP ZAP, and a release gate that blocks deployments with critical vulnerabilities. Post-implementation analysis demonstrates robust mitigation of the prioritized vulnerabilities, enhanced session integrity, and improved client-side defenses, significantly elevating the application's resilience to common web-based attacks while preserving development agility.

I. INTRODUCTION

The rapid proliferation of web-based applications has fundamentally transformed the cybersecurity landscape, shifting much of the attack surface to the application layer. In practice, modern web applications frequently fail not due to weak cryptography but because of logical flaws in access control and request integrity. Even a simple CRUD system, like the Flask To-Do application used in this project, can expose critical vulnerabilities when object-level authorization is absent or improperly enforced, or when browser security contexts are not strictly controlled.

This project treats a baseline **Flask To-Do application** as a real pre-production system requiring **rapid, verifiable security hardening**. The work was performed as a **48-hour DevSecOps sprint** for the CYC386 course, simulating a high-pressure environment typical of agile software security operations.

The target system is a **session-based web application** that supports user registration, login/logout, task creation, task listing, task editing, and task deletion. A key aspect of the application is the strict ownership of data: each task record is tied to a

specific user via a `user_id` foreign key. Therefore, the main security goal extends beyond verifying **who the user is** (authentication) to ensuring **what data the user can access or modify** (granular authorization).

Based on the application's architecture, the security remediation focused on three primary areas:

- **IDOR (Insecure Direct Object Reference):** Ensuring that each object operation (view, edit, delete) validates ownership to prevent unauthorized access to other users' tasks.
- **CSRF (Cross-Site Request Forgery):** Enforcing anti-CSRF tokens on all state-changing requests to guarantee request integrity and origin authenticity.
- **Clickjacking:** Implementing restrictive HTTP headers (X-Frame-Options, Content-Security-Policy) to prevent UI redress attacks through framing.

Beyond direct code-level mitigations, the sprint incorporated **DevSecOps best practices** through a fully automated **CI/CD pipeline using GitHub Actions**. The pipeline implemented a **shift-left security approach**, running Static Application Security Testing (SAST) via CodeQL, Dynamic Application Security Testing (DAST) with OWASP ZAP, and automated security-specific unit tests on every commit. Release gating ensured that code introducing critical vulnerabilities could not be merged, thereby maintaining a consistently hardened and secure application.

A. Project Objectives

The primary objectives of this security sprint are formally defined as follows:

1. Conduct a structured asset-based Protection Needs Elicitation (PNE) to formally identify and classify all assets requiring protection within the Flask To-Do application ecosystem.
2. Apply STRIDE threat modeling methodology to decompose the application's attack surface systematically and derive measurable, testable security requirements.
3. Quantify identified vulnerabilities using the CVSS v3.1 scoring framework to establish a data-driven remediation priority order.
4. Implement code-level security controls addressing IDOR, CSRF, and Clickjacking within a 48-hour sprint constraint.
5. Construct and validate a multi-stage DevSecOps CI/CD pipeline integrating SAST, DAST, unit testing, and critical vulnerability release gating.

B. Document Organization

The remainder of this report is structured as follows: Section II presents the Protection Needs Elicitation findings. Section III details the STRIDE threat model and attack surface analysis. Section IV presents the CVSS v3.1 risk assessment. Section V describes the identified vulnerability findings. Section VI details the code-level security implementation. Section VII covers the CI/CD pipeline architecture. Section VIII presents SAST and DAST validation results. Section IX concludes with future hardening recommendations.

II. PROTECTION NEEDS ELICITATION (PNE)

Protection Needs Elicitation (PNE) forms the foundational step of the secure software development lifecycle. Before implementing technical controls, it is imperative to systematically identify the assets requiring protection, understand the threat context surrounding those assets, and derive explicit, testable security requirements. PNE bridges the gap between high-level business requirements and concrete engineering security specifications.

A. Asset Identification and Classification

Assets within the scope of the Flask To-Do application were identified and categorized based on their role within the system architecture. For each asset, the specific Protection Needs were mapped to the classic CIA (Confidentiality, Integrity, Availability) triad.

TABLE I: Assets and Protection Needs

Asset	Components	CIA Classification	Security Rationale
User Accounts	username, password_hash, session token	Confidentiality, Integrity	Compromise enables full account takeover
Task Data	title, status, created_at, user_id	Confidentiality, Integrity, Availability	Cross-user exposure violates privacy and data isolation
Database (todo.db)	SQLite persistent store	Integrity, Availability	SQL injection or file access compromises all records
Session Cookies	Browser-managed authentication tokens	Confidentiality, Integrity	Session hijacking enables unauthorized access
CI/CD Workflow	GitHub Actions pipeline, build scripts	Integrity, Traceability	Pipeline tampering undermines security validation

B. Threat Context Analysis

The threat context for this application is defined by its architectural characteristics. The application possesses multiple state-changing endpoints (POST routes for creation, editing, deletion). Crucially, these endpoints utilize predictable object identifiers — sequential integer task_ids — exposed directly in the URL path (e.g., /tasks/3/edit). Furthermore, the application relies on form-based POST actions executed within a browser context, rendering it susceptible to cross-origin manipulation if anti-CSRF mechanisms are absent. Finally, because the UI is rendered in HTML and executed within a browser — which may navigate to hostile web contexts — the application faces inherent risks from UI redress attacks.

C. Elicited Security Requirements

Based on the asset analysis and threat context, the following six explicit security requirements were elicited to guide the implementation phase:

ID	Security Requirement
SR-1	Enforce object-level authorization on every task operation by strictly binding database queries to the authenticated user's identity.
SR-2	Protect all unsafe HTTP methods (POST, PUT, DELETE) with cryptographic CSRF tokens validated server-side prior to processing.
SR-3	Inject anti-clickjacking HTTP response headers (X-Frame-Options and Content-Security-Policy) on all server responses.
SR-4	Harden session cookie attributes, specifically enforcing the SameSite attribute to mitigate cross-site request flooding.
SR-5	Integrate automated Static Application Security Testing (SAST) and Dynamic Application Security Testing (DAST) into the SDLC.
SR-6	Implement a CI/CD release gate that automatically fails the pipeline if critical or high-severity security findings are detected.

III. THREAT MODELING AND ATTACK SURFACE ANALYSIS

To systematically analyze the attack surface, the STRIDE (Spoofing, Tampering, Repudiation, Information Disclosure, Denial of Service, Elevation of Privilege) methodology was applied. STRIDE, originally developed at Microsoft, provides a comprehensive categorization of threats, ensuring that no fundamental security property is overlooked during the design and review phases of the software development lifecycle.


A. STRIDE Summary Analysis

TABLE II: STRIDE Threat Matrix

STRIDE Category	App-Specific Threat	Impact	Current Status	Mitigation Strategy
Spoofing (S)	Unauthorized login / session hijacking	Account compromise	Mitigated	PBKDF2 password hashing; @login_required decorator
Tampering (T)	IDOR via task_id manipulation in URL	Data integrity loss	Mitigated	User-scoped DB queries (id + user_id filter)
Repudiation (R)	User denies performing destructive action	Forensic gap	Residual Gap	Future: structured audit logging required
Information Disclosure (I)	Cross-user task enumeration; debug mode stack traces	Privacy breach	Mitigated	Per-user query filtering; debug mode disabled
Denial of Service (D)	Endpoint flooding of /login (password hashing cost)	Service unavailability	Residual Gap	Future: Flask-Limiter rate limiting middleware
Elevation of Privilege (E)	Horizontal privilege escalation to foreign data	Access boundary bypass	Mitigated	Object-level authorization enforcement

B. Data Flow Diagram (DFD) Trust Boundary Analysis

The following DFD describes the data flows between the browser, the Flask application, the ORM layer, and the SQLite database, including the security controls applied at each layer.



The system architecture is structured into three distinct trust boundaries that act as security gates for data as it moves from the user to the database:

TB1 — Untrusted Input Boundary (Client Layer)

This boundary separates the Browser from the server. All data coming from the user agent is considered hostile or untrusted. Security at this transition focuses on verifying that the request is legitimate (Authentication) and that it was intentionally sent by the user (CSRF verification).

TB2 — Flask App / Security Layer (Application Layer)

This is the internal processing zone where the Flask Blueprints reside. Once a request crosses TB1, it enters this layer to be filtered. Here, middleware enforces access rules (like `@login_required`) and injects security headers (X-Frame-Options) into the response to protect the user from UI-based attacks like clickjacking.


TB3 — Trusted App / Data Boundary (Data Layer)

This boundary sits between the application logic and the SQLite database. It is the most critical line of defense for data privacy. At this stage, the SQLAlchemy ORM applies Ownership Checks (IDOR mitigation). It ensures that even if a request is authenticated, the user only has permission to read or write data where the `user_id` matches their own session identity.

C. Attack Tree Decomposition

An attack tree was constructed to visually and logically map the adversarial paths leading to the primary goal: 'Unauthorized Task Access or Modification.' The tree utilizes an OR node at the root, indicating that an attacker needs only one successful path to achieve their goal. Three primary attack paths were identified and analyzed:

Attack Path	Description	Outcome	Defense Mechanism
Path 1: IDOR	Authenticated user manipulates integer <code>task_id</code> in URL to access foreign record	BLOCKED	User-scoped DB query returns HTTP 404 — prevents enumeration of foreign IDs
Path 2: CSRF	Attacker hosts page with hidden form auto-submitting POST to victim's authenticated session	BLOCKED	Flask-WTF rejects request with HTTP 400 — missing session-bound CSRF token
Path 3: Clickjacking	Attacker frames legitimate Flask app in transparent iframe to hijack victim clicks	BLOCKED	X-Frame-Options: DENY and CSP frame-ancestors 'none' prevent browser rendering in frames



IV. CVSS v3.1 RISK ASSESSMENT

To quantify the severity of the identified vulnerabilities and justify the prioritization of the 48-hour sprint, the Common Vulnerability Scoring System (CVSS) version 3.1 was utilized. CVSS provides a standardized, vendor-neutral framework for communicating the characteristics and severity of software vulnerabilities, enabling objective, reproducible risk scoring.

TABLE III: CVSS v3.1 Risk Ratings for Identified Vulnerabilities

Vulnerability	Base Score	Severity	CVSS Vector	Key Rationale
IDOR	6.5	Medium	AV:N/AC:L/PR:L/UI:N/S:U/C:H/I:N/A:N	Network-accessible; low complexity; requires valid account (PR:L); high confidentiality impact
CSRF	8.8	High	AV:N/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H	No privileges required on target; victim-assisted; full C/I/A impact on worst-case state-change
Clickjacking	5.4	Medium	AV:N/AC:L/PR:N/UI:R/S:U/C:L/I:L/A:N	Requires user interaction; limited impact — attacker can only force victim to perform authorized actions

A. IDOR — Detailed CVSS Analysis (Base Score: 6.5 Medium)

The IDOR vulnerability receives a Medium severity rating of 6.5. The Attack Vector is Network (AV:N), as the manipulation occurs over HTTP. Attack Complexity is Low (AC:L), requiring no specialized conditions. The vulnerability requires Low Privileges (PR:L) — an attacker must possess a valid user account to be authenticated. No User Interaction is required (UI:N) beyond the initial authentication. Scope is Unchanged (S:U). The Confidentiality impact is High (C:H) for read-based IDOR, as an attacker can access foreign task titles and metadata. Integrity and Availability impacts depend on whether write operations are also vulnerable. These scores directly justified IDOR's placement as a critical code-level remediation priority.

B. CSRF — Detailed CVSS Analysis (Base Score: 8.8 High)

CSRF receives a High severity rating of 8.8, reflecting its victim-assisted but highly impactful nature. The Attack Vector is Network (AV:N). No Privileges are required on the target system (PR:N) because the attack leverages the victim's existing authenticated session. User Interaction (UI:R) is required — the victim must visit the attacker's hosted page. Because a successful CSRF attack can trigger state-changing operations such as deleting all tasks, creating spurious records, or modifying task content, the impact on Confidentiality, Integrity, and Availability are all rated High (C:H/I:H/A:H) in the worst-case scenario. This score designated CSRF as the highest-priority fix in the 48-hour sprint timeline.

C. Clickjacking — Detailed CVSS Analysis (Base Score: 5.4 Medium)

Clickjacking receives a Medium severity rating of 5.4. While the attack vector is Network (AV:N) with Low complexity (AC:L), the strict requirement for User Interaction (UI:R) substantially limits its exploitability. The attack constrains the attacker to forcing the victim to perform actions they are already authorized to perform — the impact is limited (C:L/I:L) and Availability is not directly impacted (A:N). Despite the Medium rating, Clickjacking was addressed within the sprint as the mitigation is trivially inexpensive (a single after_request header hook) relative to the risk elimination value it provides.

V. VULNERABILITY FINDINGS AND REMEDIATION REQUIREMENTS

A. Finding 1: Insecure Direct Object Reference (IDOR)

Description: The application utilized direct object references in its URL routing scheme (e.g., /tasks/<int:task_id>/edit). The vulnerability risk is that if the application logic fetched the task solely based on the task_id parameter without verifying the ownership context, an authenticated user could manipulate the integer to access or modify data belonging to another user.

Exploitation Scenario: User A creates a task assigned ID 4. User A then sends an HTTP POST request to /tasks/5/edit, a task belonging to User B. If the server executes Task.query.get(5) without a user_id ownership check, User A successfully loads and modifies User B's data. This constitutes a horizontal privilege escalation with no technical complexity beyond URL manipulation.

Security Requirement: The data access layer must explicitly enforce object-level authorization. Every database query fetching a task must include a conditional clause verifying that the task's user_id matches the currently authenticated user's ID (SR-1).

B. Finding 2: Cross-Site Request Forgery (CSRF)

Description: The application relies heavily on HTML forms utilizing the POST method for state changes. In the absence of anti-CSRF mechanisms, the browser's automatic inclusion of session cookies for cross-origin POST requests creates a severe vulnerability. The browser does not distinguish between a form submitted by the legitimate application and one submitted by an attacker's page on behalf of the victim.

Exploitation Scenario: An attacker constructs a malicious webpage containing a hidden HTML form that auto-submits a POST request to the Flask application's task deletion endpoint. When an authenticated victim visits this page, the browser automatically attaches their session cookie, and the Flask application processes the deletion as a legitimate request from the victim, resulting in unauthorized data destruction.

Security Requirement: The application must implement the Synchronizer Token Pattern (SR-2). A cryptographically strong, random token must be generated per session, embedded in all unsafe forms as a hidden field, and strictly validated on the server side before any state-altering logic is executed.

C. Finding 3: Clickjacking (UI Redress Attack)

Description: By default, web browsers allow a page to be loaded inside an HTML <iframe> element, even if the framed page originates from a different domain. This behavior allows an attacker to obscure the legitimate application's UI with deceptive overlays, causing users to unknowingly interact with the underlying application.

Exploitation Scenario: An attacker creates a page titled 'Click to Win a Prize.' Behind the visible button, the legitimate Flask app is loaded in a zero-opacity iframe, perfectly aligned so that the victim's click actually triggers the 'Delete Task' button on the hidden application. The victim perceives only the attacker's UI while performing unintended actions on the real application.

Security Requirement: The server must instruct the browser to refuse framing of its content (SR-3). This is achieved by deploying the X-Frame-Options HTTP header set to 'DENY' and the modern Content-Security-Policy directive 'frame-ancestors none' to cover both legacy and modern browser contexts.

VI. SECURITY IMPLEMENTATION (CODE-LEVEL)

The translation of security requirements into concrete code implementations was executed methodically across the application's routing, templating, and configuration layers. Each implementation was designed for reusability, centralization, and testability to prevent future bypass through code evolution.

A. IDOR Mitigation — User-Scoped Object Retrieval

The IDOR vulnerability was resolved by abstracting the task retrieval logic into a dedicated helper function within `app/routes/tasks.py`. This centralization ensures that the authorization check cannot be accidentally bypassed by individual route handlers and provides a single point of maintenance.

```
# app/routes/tasks.py
from flask_login import current_user
from app.models import Task

def _get_user_task_or_404(task_id):
    """
    Retrieves a task ONLY if it belongs to the current user.
    Returns 404 if the task does not exist OR belongs to another user.
    This prevents IDOR by unifying object existence and authorization.
    """
    return Task.query.filter_by(
        id=task_id,
        user_id=current_user.id
    ).first_or_404()
```

Architectural Impact: By utilizing SQLAlchemy's `filter_by` method to join both the `id` and `user_id` columns in a single database query, the authorization check is performed atomically at the data layer. If a user queries a foreign ID, the database returns an empty result set, and `first_or_404()` triggers an HTTP 404 Not Found response. This approach provides a dual benefit: it mitigates IDOR and prevents attackers from using the application to enumerate the existence of foreign task IDs — returning 404 instead of 403 avoids revealing whether an ID is valid but unauthorized.

B. CSRF Mitigation — Synchronizer Token Pattern

CSRF protection was implemented globally using the Flask-WTF extension, providing session-bound token generation and validation across all state-changing endpoints.

```
# app/__init__.py — Global CSRF initialization
from flask_wtf.csrf import CSRFProtect

csrf = CSRFProtect()

def create_app(config):
    app = Flask(__name__)
    app.config.from_object(config)
    csrf.init_app(app) # Registers CSRF protection globally
    return app

<!-- Template integration — all unsafe forms (Jinja2) -->
<form method="POST" action="/tasks/create">
    <input type="hidden" name="csrf_token" value="{{ csrf_token() }}">
    <!-- form fields -->
</form>
```

Mechanism: Flask-WTF automatically intercepts all incoming POST, PUT, PATCH, and DELETE requests. It compares the submitted `csrf_token` against the securely stored, session-bound token (generated using `os.urandom` and stored in the server-side session). If the tokens mismatch or are absent, the framework immediately aborts the request with an HTTP 400 Bad Request error. Since the attacker's forged page cannot read the victim's session token (enforced by the browser's Same-Origin Policy), the forged request will always lack a valid token, rendering the CSRF vector inert.

C. Session Hardening and Clickjacking Prevention

Session hardening addresses CSRF at the browser level as a defense-in-depth measure, while the response header injection directly eliminates the Clickjacking attack surface.

```
# config.py — Session cookie hardening
SESSION_COOKIE_SAMESITE = "Lax"
SESSION_COOKIE_HTTPONLY = True
SESSION_COOKIE_SECURE = True # Enforce in production with HTTPS
```

```
# app/__init__.py — Universal security header injection
@app.after_request
def add_security_headers(response):
    response.headers["X-Frame-Options"] = "DENY"
    response.headers["Content-Security-Policy"] = "frame-ancestors 'none'"
    response.headers["X-Content-Type-Options"] = "nosniff"
    response.headers["Referrer-Policy"] = "strict-origin-when-cross-origin"
    return response
```

Impact: The SameSite=Lax cookie attribute instructs modern browsers to withhold the session cookie on cross-site POST requests, neutralizing most CSRF attacks at the browser level as a complementary defense. The after_request decorator ensures security headers are unconditionally attached to every HTTP response. X-Frame-Options: DENY maintains legacy browser compatibility, while frame-ancestors 'none' provides stricter, modern enforcement. Additional headers (X-Content-Type-Options and Referrer-Policy) were included to further reduce the browser-level attack surface.

D. Automated Security Test Suite

To validate the implemented controls and ensure continuous verification within the CI/CD pipeline, a dedicated security test suite was developed in tests/test_security.py. The suite contains the following five test cases:

#	Test Case	Description	Expected Response	Requirement Validated
1	Unauthenticated Redirect	Access protected route without session	HTTP 302 → /login	SR-1, Authentication
2	Owner-Only Task Visibility	Task list returns only authenticated user's tasks	HTTP 200, filtered results	SR-1 (IDOR)
3	IDOR Block on Foreign Task	Edit page request for foreign task_id	HTTP 404	SR-1 (IDOR)
4	CSRF Rejection on Task Create	POST to /tasks/create without csrf_token	HTTP 400	SR-2 (CSRF)
5	CSRF Rejection on Logout	POST to /logout without csrf_token	HTTP 400	SR-2 (CSRF)


VII. CI/CD PIPELINE INTEGRATION AND DEVSECOPS AUTOMATION

To ensure the sustainability of the security hardening and prevent the introduction of future regression vulnerabilities, a comprehensive DevSecOps pipeline was engineered using GitHub Actions. This pipeline operationalizes the 'shift-left' security principle by embedding automated security testing and quality gates directly into the development workflow, providing immediate feedback to developers before code can be merged to production branches.

A. Trigger Strategy

The pipeline is configured to execute automatically on two distinct trigger events:

- **push:** Triggered when code is pushed directly to the main or development branches, ensuring immediate validation of all committed changes.
- **pull_request:** Triggered when a pull request is opened or updated against the target branch. This is critical for the shift-left paradigm, providing security feedback to developers and code reviewers before any code is merged to the protected main branch.



Active GitHub Pull Request demonstrating the DevSecOps workflow. The interface shows the feature-to-main branch transition, peer reviewer assignment, and the status of integrated CI security checks.

B. Pipeline Job Architecture

The workflow is decomposed into four sequential, dependent jobs that collectively form a comprehensive security quality gate. Each job must succeed before the next is initiated.

Job	Name	Description	Failure Condition
1	Test	Installs dependencies and executes all security unit tests via pytest	Any test case fails
2	SAST (CodeQL)	Performs semantic code analysis; traces tainted data flows from HTTP inputs to DB sinks	Critical/High code scanning alert raised
3	DAST (ZAP)	Executes OWASP ZAP baseline scan against running app; validates HTTP security headers	High-severity passive/active alert found
4	Security Gate	Queries GitHub Code Scanning API for Critical/High alerts on current commit ref	Any critical/high alerts exist on commit

C. GitHub Actions Workflow Configuration

```
# .github/workflows/ci-cd.yml
name: Flask CI/CD Security Pipeline
on:
  push:
    branches: [main, development]
  pull_request:
    branches: [main]

jobs:
  build-test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - name: Setup Python 3.13
        uses: actions/setup-python@v4
        with:
          python-version: '3.13'
      - name: Install Dependencies
        run: pip install -r requirements.txt
      - name: Run Security Test Suite
        run: pytest tests/test_security.py -v


  codeql-sast:
    needs: build-test
    uses: github/codeql-action/init@v3
    with:
      languages: python

  owasp-dast:
    needs: codeql-sast
    steps:
      - name: Start Flask Application
        run: python run.py &
      - name: ZAP Baseline Scan
        uses: zaproxy/action-baseline@v0.11.0
        with:
          target: 'http://localhost:5000'
          artifact_name: zap-baseline-report

  security-gate:
    needs: owasp-dast
    steps:
      - name: Check for Critical Alerts
        run: |
          ALERTS=$(gh api /repos/$REPO/code-scanning/alerts \
            --jq '[.[] | select(.state=="open")] | length')
          if [ $ALERTS -gt 0 ]; then exit 1; fi
```

D. Artifact Generation and Evidence Retention

The ZAP DAST job is configured to upload the zap-baseline-report as a GitHub Actions artifact. This enables developers and security reviewers to download the complete HTML and JSON reports, inspect detailed pass/fail metrics for each scanned rule, and track the remediation of DAST findings across successive pipeline runs. Artifacts are retained for 30 days per pipeline execution, providing a time-indexed audit trail of the application's security posture evolution.



Visual representation of the GitHub Actions CI/CD pipeline showing interdependent security jobs and successful artifact generation.


VIII. SAST + DAST RESULTS AND VALIDATION

The implementation of the security controls was validated through a before-and-after comparative analysis utilizing both SAST and DAST methodologies. This comparative approach provides verifiable, tool-generated evidence of vulnerability remediation that constitutes a formal security acceptance criterion.


A. SAST Results — CodeQL Analysis

Pre-Fix State: Initial conceptual CodeQL analysis highlighted weaknesses in the application's access control logic. The semantic analysis would flag areas where database queries relied on unsanitized, user-provided identifiers (the `task_id` route parameter) without any verification against the authenticated session context. CodeQL's data-flow analysis treats the `task_id` URL parameter as a tainted source; without the `user_id` filter, this tainted value flows unchecked to the database query sink.

Post-Fix State: Following implementation of the `_get_user_task_or_404()` helper function, subsequent CodeQL scans clear previous alerts. The tool's data-flow analysis confirms that user input (`task_id` from the route) is now formally bound to the authenticated session context (`current_user.id`) before reaching the database sink. The tainted data is sanitized by the ownership check, satisfying CodeQL's security query requirements. The critical vulnerability gate remains green, permitting pipeline progression.



GitHub CodeQL dashboard identifying a high-severity "Flask debug mode" finding during static analysis.



GitHub repository security overview confirming active status for Code Scanning and Secret Scanning alerts.

B. DAST Results — OWASP ZAP Baseline Scan

Pre-Fix Theoretical Baseline: A ZAP baseline scan against the un-hardened application would flag high-confidence alerts including: (1) Anti-Clickjacking Header missing — X-Frame-Options not set; (2) Content-Security-Policy Header missing; (3) Session cookie without SameSite attribute in legacy browser contexts.

Post-Fix Results: Post-implementation ZAP baseline reports demonstrate complete compliance with all scanned rules. The HTTP response analyzer confirms all required headers are present and correctly configured, with no critical or high-severity passive or active alerts generated during automated scanning.

TABLE IV: ZAP Baseline Scan Results — Pre and Post Remediation

ZAP Rule	Pre-Fix Status	Post-Fix Status	Validation Evidence
Anti-Clickjacking Header	FAIL	PASS	X-Frame-Options: DENY present in all responses
Content-Security-Policy	FAIL	PASS	frame-ancestors 'none' verified in CSP header
Session Cookie SameSite	FAIL	PASS	SameSite=Lax confirmed in Set-Cookie header
CSRF Token Validation	NOT TESTED	PASS	Unit tests confirm HTTP 400 on missing tokens
IDOR Access Control	NOT TESTED	PASS	Unit tests confirm HTTP 404 on foreign task_id

C. Unit Test Execution Evidence

Local continuous testing during the development phase, mirrored by the CI pipeline execution, confirmed that all five security-specific tests within tests/test_security.py pass consistently across all code revisions. The tests validate the defensive boundaries of the implemented controls: unauthenticated access is redirected (HTTP 302), foreign object access is denied (HTTP 404), and missing CSRF tokens are rejected (HTTP 400). All five tests passing constitutes formal evidence of requirements SR-1 through SR-3 being met.

```
PS C:\Users\mahee\OneDrive\Documents\flask_todo_app> python -m unittest discover -s tests -p "test*.py" -v
test_create_task_without_csrf_token_is_rejected (test_security.CsrfEnforcementTests.test_create_task_without_csrf_token_is_rejected) ...
ok
test_authenticated_user_sees_only_own_tasks (test_security.SecurityTests.test_authenticated_user_sees_only_own_tasks) ... C:\Python314\Li
b\site-packages\sqlalchemy\sql\schem.py:3624: DeprecationWarning: datetime.datetime.utcnow() is deprecated and scheduled for removal in
a future version. Use timezone-aware objects to represent datetimes in UTC: datetime.datetime.now(datetime.UTC).
    return util.wrap_callable(lambda ctx: fn(), fn) # type: ignore
ok
test_idor_block_foreign_task_edit (test_security.SecurityTests.test_idor_block_foreign_task_edit) ... ok
test_unauthenticated_redirect_from_index (test_security.SecurityTests.test_unauthenticated_redirect_from_index) ... ok
-----
OK
```

Terminal execution log of the unittest discovery suite. The output confirms that all security-specific test cases—including CSRF enforcement and IDOR mitigation—are functioning correctly and passing within the local development environment

```

PS C:\Users\mahee\OneDrive\Documents\flask_todo_app> python -m pytest -v
===== test session starts =====
platform win32 -- Python 3.14.2, pytest-9.0.3, pluggy-1.6.0 -- C:\Python314\python.exe
cachedir: .pytest_cache
rootdir: C:\Users\mahee\OneDrive\Documents\flask_todo_app
collected 5 items

tests/test_security.py::SecurityTests::test_authenticated_user_sees_only_own_tasks PASSED [ 20%]
tests/test_security.py::SecurityTests::test_idor_block_foreign_task_edit PASSED [ 40%]
tests/test_security.py::SecurityTests::test_unauthenticated_redirect_from_index PASSED [ 60%]
tests/test_security.py::CsrfEnforcementTests::test_create_task_without_csrf_token_is_rejected PASSED [ 80%]
tests/test_security.py::CsrfEnforcementTests::test_logout_without_csrf_token_is_rejected PASSED [100%]

===== 5 passed in 2.16s =====
PS C:\Users\mahee\OneDrive\Documents\flask_todo_app>

```

Terminal output showing successful execution of the security unit test suite. All 5 critical security tests (IDOR, CSRF, and Authentication) passed in 2.16s

IX. CONCLUSION AND FUTURE HARDENING RECOMMENDATIONS

A. Summary of Achievements

This 48-hour DevSecOps security sprint successfully achieved all primary security objectives established for the Flask To-Do application. The project demonstrated a methodical, standards-driven approach to application security, transitioning from theoretical vulnerability analysis to practical, verified code-level remediations. The core achievements of this sprint are:

- Complete mitigation of Insecure Direct Object Reference (IDOR) through the implementation of strict, user-scoped database object queries via the `_get_user_task_or_404()` helper, fundamentally securing the horizontal authorization boundary.
- Robust neutralization of Cross-Site Request Forgery (CSRF) risks via the Flask-WTF Synchronizer Token Pattern, coupled with session cookie `SameSite=Lax` hardening as defense-in-depth.
- Prevention of Clickjacking attacks through systematic injection of X-Frame-Options: DENY and Content-Security-Policy frame-ancestors 'none' headers via the `after_request` hook.
- Successful operationalization of DevSecOps principles through a multi-stage GitHub Actions CI/CD pipeline integrating SAST (CodeQL), DAST (OWASP ZAP), automated security unit tests, and a critical-failure release gate.

The resulting system is demonstrably more attack-resistant than its baseline iteration. All three targeted OWASP Top 10 vulnerability categories have been addressed with verified, automated testing evidence. The CI/CD pipeline ensures that future development activity cannot regress these security properties without explicit pipeline failure.

B. Residual Risks and Future Hardening Roadmap

Despite the significant security improvements achieved, the threat modeling process identified residual gaps that exceed the scope of a 48-hour sprint but must be addressed prior to true production deployment:

Priority	Recommendation	Rationale	Addresses STRIDE
P1	Rate Limiting (Flask-Limiter)	Prevent credential stuffing and DoS on /login and /register endpoints	Denial of Service (D)
P2	Structured Audit Logging	Append-only log of auth events and data mutations with user ID, IP, and timestamp	Repudiation (R)
P3	Enforce debug=False in Production	Prevent stack trace and environment variable disclosure via Flask debugger	Information Disclosure (I)

P4	Multi-Factor Authentication (MFA)	Reduce account compromise risk for high-privilege operations	Spoofing (S)
P5	Authenticated DAST Scanning	Configure ZAP with session context to scan behind the login wall and test IDOR dynamically	Tampering (T)
P6	Database Encryption at Rest	Adopt TDE-capable database for production; protect user data if file-system access is obtained	Information Disclosure (I)
P7	HSTS and Extended Security Headers	Add Strict-Transport-Security, X-Content-Type-Options, and Referrer-Policy headers	Information Disclosure (I)

C. DevSecOps Maturity Assessment

The completion of this sprint advances the application from DevSecOps Maturity Level 0 (no security automation) to Level 2 (integrated automated security testing with quality gates). Progression to Level 3 (continuous security monitoring with threat intelligence integration) requires implementation of the roadmap items above, specifically structured audit logging with SIEM integration, periodic automated penetration testing, and dependency vulnerability scanning via Dependabot or Snyk.


X. ACKNOWLEDGMENTS

The author acknowledges Engr. Muhammad Ahmad Nawaz for guidance and mentorship throughout the CYC386 Secure Software Design and Development course. Appreciation is extended to peers Maheen Aamir , Malik Muhammad Huzaifa , Muzamil Ali Akhtar and Maryum Masood for collaborative discussions and technical insights during the 48-hour sprint. The DevSecOps methodology applied in this project draws from OWASP guidance, industry best practices, and academic literature in secure software engineering.


APPENDICES

Appendix A — Evidence Checklist

#	Evidence Item	Required Format	Status
1	Pull Request created in GitHub with reviewer assigned	Screenshot	[]
2	GitHub Actions workflow showing all 4 jobs passing	Screenshot	[]
3	GitHub CodeQL results showing zero critical/high alerts	Screenshot	[]
4	OWASP ZAP artifact downloaded from GitHub Actions run	ZIP/HTML Report	[]
5	THREAT_MODEL.pdf in root of repository	PDF File	[]
6	Final Report.pdf attached to submission	PDF File	[]
7	Demo video (5–7 minutes) demonstrating application, mitigations, and pipeline	Video Link	[]



GitHub Pull Request (PR) interface showing the transition of security documentation and code hardening from the feature branch to the main branch, including assigned peer reviewer for manual audit.



GitHub Actions dashboard showing successful SAST, DAST, and security gate completion.

Appendix B — Full STRIDE Threat Matrix (Extended)

Detailed threat decomposition for each STRIDE category mapped to specific application endpoints, including attack conditions, threat actors, and complete mitigation mappings. Refer to Section III and Table II for the core STRIDE analysis. Full tabular detail available in the supplementary THREAT_MODEL.pdf repository artifact.

Appendix C — CVSS v3.1 Scoring Breakdown

Complete breakdown of CVSS v3.1 base metric group scores for each identified vulnerability, including Base Score, Temporal Score adjustments (where applicable), and Environmental Score considerations for the target deployment context. Refer to Section IV and Table III for primary scoring data.

Vulnerability	AV	AC	PR	UI	S	C	I	A
IDOR	N	L	L	N	U	H	N	N
CSRF	N	L	N	R	U	H	H	H
Clickjacking	N	L	N	R	U	L	L	N

AV=Attack Vector, AC=Attack Complexity, PR=Privileges Required, UI=User Interaction, S=Scope, C=Confidentiality, I=Integrity, A=Availability | N=Network/None, L=Low, R=Required, H=High, U=Unchanged

Appendix D — CI/CD Workflow YAML (Full)

Full GitHub Actions configuration file for automated testing, SAST, DAST, and critical vulnerability gating. See Section VII-C for the annotated workflow configuration. The complete file is located at `.github/workflows/ci-cd.yml` in the project repository.

Appendix E — Security Unit Test Suite (`test_security.py`)

Complete test file with all five test cases, setup fixtures, and teardown logic. Each test validates a specific security requirement (SR-1 through SR-3) as defined in Section II-C. Test descriptions, expected HTTP response codes, and coverage metrics are detailed in Section VI-D and Table V above.

Appendix F — Data Flow Diagram

Layered visualization of client, application, ORM, and database trust boundaries, including data flows annotated with trust levels, control enforcement points, and vulnerability injection nodes. The DFD was used to identify the trust boundary transitions analyzed in Section III-B.